

A large orange graphic consisting of a wavy line that ends in a right-pointing arrow.

## **pg\_stat\_advisor - optimizing PostgreSQL planner statistics with recommendations**

Ilia Evdokimov, Tantor Labs LLC

# \* Contents



- issues with statistics
- **pg\_stat\_advisor** overview
- future plans

## \* Compare estimated and actual rows

For the best plan **estimated rows** should be similar with **actual rows**:

```
EXPLAIN ANALYZE
SELECT * FROM my_tbl
WHERE fld_1 = 500 AND fld_2 = 100;
>>>
    Gather  (cost=1000.00..124002.73 rows=1 width=12) (actual
time=5.778..135.219 rows=100 loops=1)"
    Workers Planned: 2"
    Workers Launched: 2"
    -> Parallel Seq Scan on my_tbl  (cost=0.00..122975.12 rows=115
width=12) (actual time=86.647..129.191 rows=33 loops=3)"
        Filter: ((fld_1 = 500) AND (fld_2 = 100))"
        Rows Removed by Filter: 333300"
    Planning Time: 0.160 ms"
    Execution Time: 135.260 ms"
```

## \* Prepare test table



Create table with disabled `autovacuum` to update statistics manually:

```
CREATE TABLE my_tbl(fld_1 INTEGER, fld_2 BIGINT)
WITH (autovacuum_enabled = false);
>>>
CREATE TABLE
```

## \* Evaluate data cardinality



Evaluate test data before **INSERT**:

```
SELECT
    COUNT(DISTINCT T.fld_1)::float/10000000,
    COUNT(DISTINCT T.fld_2)::float/10000000
FROM (
    SELECT i/100 AS fld_1, i/500 AS fld_2
    FROM generate_series(1, 10000000) s(i)
) T;
>>>
    0.01      0.002      <-- low cardinality
```

## \* Insert values

Fill the table with diverse data:

```
INSERT INTO my_tbl (fld_1, fld_2)
SELECT
    i/100 as fld_1,
    i/500 as fld_2
FROM generate_series(1, 10000000) s(i);
>>>
INSERT 0 1000000
```

# \* Compare estimated and actual rows

Without **ANALYZE**

```
EXPLAIN ANALYZE
SELECT * FROM my_tbl
WHERE fld_1 = 500 AND fld_2 = 100;
>>>
    Gather  (cost=1000.00..124002.73 rows=276 width=12) (actual
time=5.778..135.219 rows=100 loops=1)"
    Workers Planned: 2"
    Workers Launched: 2"
    -> Parallel Seq Scan on my_tbl  (cost=0.00..122975.12 rows=115
width=12) (actual time=86.647..129.191 rows=33 loops=3)"
        Filter: ((fld_1 = 500) AND (fld_2 = 100))"
        Rows Removed by Filter: 333300"
    Planning Time: 0.160 ms"
    Execution Time: 135.260 ms"
```

# \* Compare estimated and actual rows

After **ANALYZE**

```
ANALYZE my_tbl;
>>>
ANALYZE

EXPLAIN ANALYZE
SELECT * FROM my_tbl
WHERE fld_1 = 500 AND fld_2 = 100;
>>>
    Gather  (cost=1000.00..117554.23 rows=1 width=12) (actual
time=6.061..134.014 rows=100 loops=1)"
    Workers Planned: 2"
    Workers Launched: 2"
    -> Parallel Seq Scan on my_tbl  (cost=0.00..116554.12 rows=1
width=12) (actual time=86.053..128.202 rows=33 loops=3)"
        Filter: ((fld_1 = 500) AND (fld_2 = 100))"
        Rows Removed by Filter: 333300"
    Planning Time: 0.128 ms"
    Execution Time: 134.057 ms"
```



## \* default\_statistics\_target

```
SHOW default_statistics_target;
>>>
    default_statistics_target
-----
    100                      <-- this is default value
(1 row)

SELECT reltuples FROM pg_class WHERE relname = 'my_tbl';
>>>
    reltuples
-----
  100000000
(1 row)
```

## \* Compare estimated and actual rows



After `default_statistics_target = 1000`

```
SHOW default_statistics_target;
```

```
>>>
```

```
1000
```

```
EXPLAIN ANALYZE
```

```
SELECT * FROM my_tbl
```

```
WHERE fld_1 = 500 AND fld_2 = 100;
```

```
>>>
```

```
      Gather  (cost=1000.00..117554.23 rows=1 width=12) (actual  
time=6.061..134.014 rows=100 loops=1)"
```

```
        Workers Planned: 2"
```

```
        Workers Launched: 2"
```

```
      -> Parallel Seq Scan on my_tbl  (cost=0.00..116554.12 rows=1  
width=12) (actual time=86.053..128.202 rows=33 loops=3)"
```

```
        Filter: ((fld_1 = 500) AND (fld_2 = 100))"
```

```
        Rows Removed by Filter: 333300"
```

```
Planning Time: 0.128 ms"
```

```
Execution Time: 134.057 ms"
```

## \* Extended statistics



<b>dependencies</b>	Evaluates functional dependencies between columns
<b>ndistinct</b>	Counts unique combinations of values in columns
<b>mcv</b>	Analyzes the most frequent value combinations in columns

```
CREATE STATISTICS stat_tbl (dependencies)  
ON fld_1, fld_2  
FROM my_tbl;  
>>>  
CREATE STATISTICS  
  
ANALYZE my_tbl;  
>>>  
ANALYZE
```

## \* Results

```
EXPLAIN ANALYZE
```

```
SELECT * FROM my_tbl
```

```
WHERE fld_1 = 500 AND fld_2 = 100;
```

```
>>>
```

```
Gather (cost=1000.00..24324.71 rows=100 width=12) (actual  
time=102.944..114.427 rows=100 loops=1)"
```

```
Workers Planned: 2"
```

```
Workers Launched: 2"
```

```
-> Parallel Seq Scan on my_tbl (cost=0.00..23311.01 rows=1 width=12)  
(actual time=68.527..101.757 rows=33 loops=3)"
```

```
Filter: ((fld_1 = 500) AND (fld_2 = 100))"
```

```
Rows Removed by Filter: 666634"
```

```
Planning Time: 0.109 ms"
```

```
Execution Time: 116.551 ms"
```

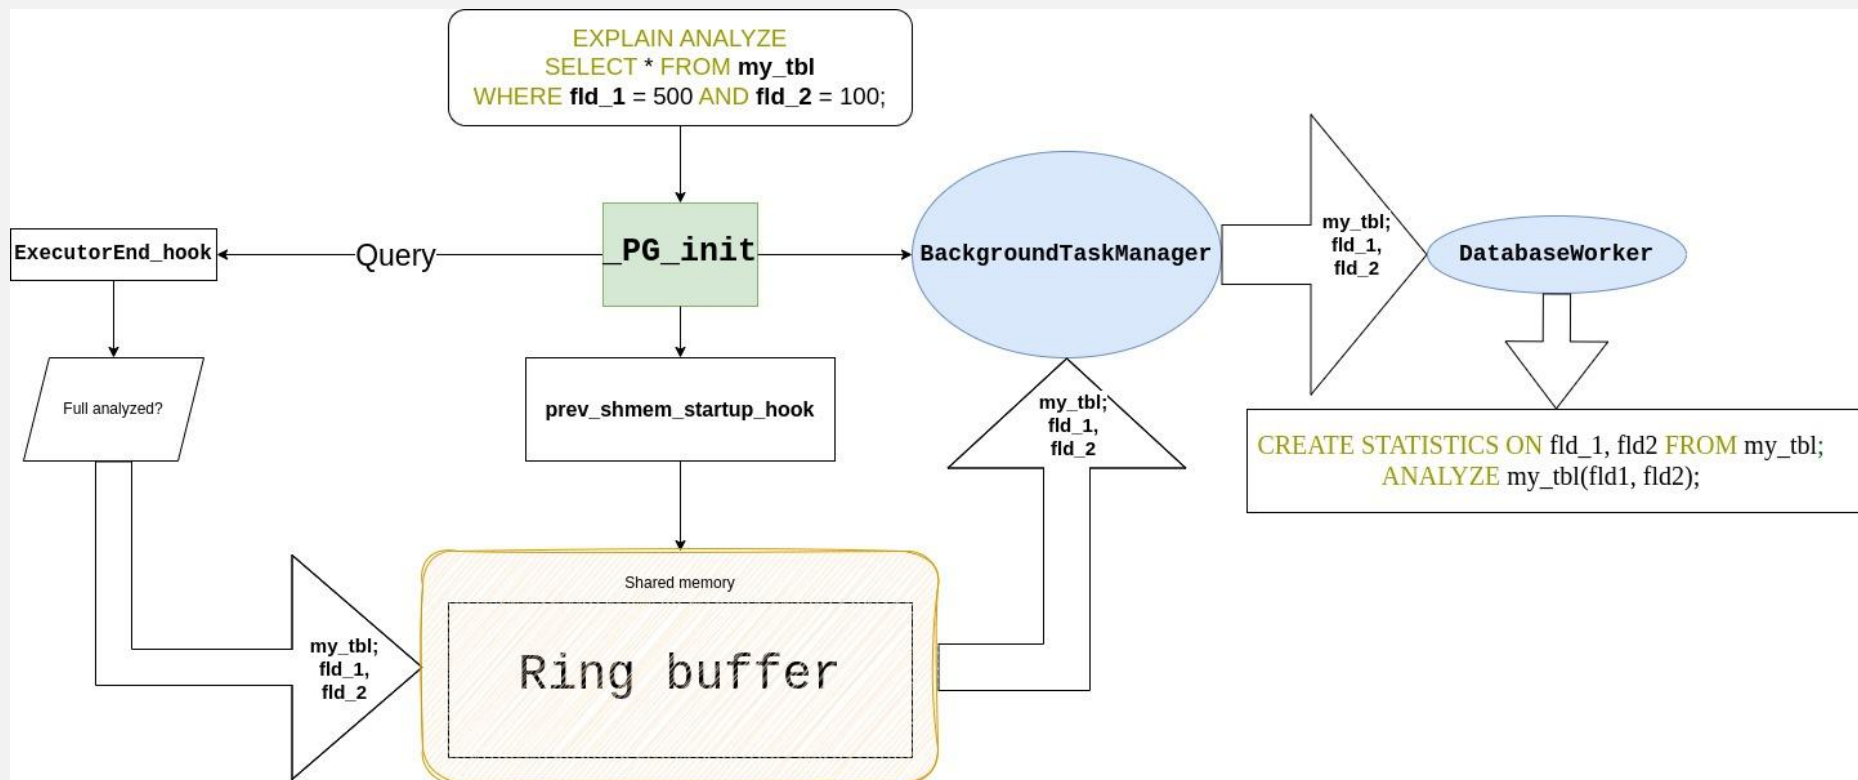
## \* About pg\_stat\_advisor



What is the Purpose of `pg_stat_advisor`?

- Suggests performing `ANALYZE` operations on non-analyzed tables.
- Creates `extended statistics and analyzes` table based on specific criteria.

# \* pg\_stat\_advisor internals



## \* How use pg\_stat\_advisor



Add library to `postgresql.conf` for current session:

```
shared_preload_libraries = 'pg_stat_advisor'
```

Set extension parameters in `postgresql.conf` or current session:

- `pg_stat_advisor.ring_buffer_capacity`
- `pg_stat_advisor.suggest_statistics_threshold`

## \* Configure ANALYZE

```
EXPLAIN ANALYZE SELECT * FROM my_tbl WHERE fld_1 = 500 AND fld_2 = 100;
>>>
NOTICE: pg_stat_advisor suggestion: 'ANALYZE my_tbl'
Gather  (cost=1000.00..124002.73 rows=276 width=12) (actual
        time=5.778..135.219 rows=100 loops=1)"
    Workers Planned: 2"
    Workers Launched: 2"
    -> Parallel Seq Scan on my_tbl (cost=0.00..122975.12 rows=115
        width=12) (actual time=86.647..129.191 rows=33 loops=3)"
        Filter: ((fld_1 = 500) AND (fld_2 = 100))"
        Rows Removed by Filter: 333300"
    Planning Time: 0.160 ms"
    Execution Time: 135.260 ms"
```



## \* Configure suggest\_statistics\_threshold



```
-- estimated_row / actual rows < pg_stat_advisor.suggest_statistics_threshold  
SET pg_stat_advisor.suggest_statistics_threshold = 0.01;
```

```
>>>
```

```
SET
```

```
EXPLAIN ANALYZE SELECT * FROM my_tbl WHERE fld_1 = 500 AND fld_2 = 100;
```

```
>>>
```

```
Gather (cost=1000.00..24311.11 rows=1 width=12) (actual  
time=104.973..113.246 rows=100 loops=1)"
```

```
Workers Planned: 2"
```

```
Workers Launched: 2"
```

```
-> Parallel Seq Scan on my_tbl (cost=0.00..23311.01 rows=1 width=12)  
(actual time=68.527..101.757 rows=33 loops=3)"
```

```
Filter: ((fld_1 = 500) AND (fld_2 = 100))"
```

```
Rows Removed by Filter: 666634"
```

```
Planning Time: 0.109 ms"
```

```
Execution Time: 116.551 ms"
```

## \* Results of applying statistics



```
EXPLAIN ANALYZE
```

```
SELECT * FROM my_tbl
```

```
WHERE fld_1 = 500 AND fld_2 = 100;
```

```
>>>
```

```
Gather (cost=1000.00..24324.71 rows=100 width=12) (actual  
time=102.944..114.427 rows=100 loops=1)"
```

```
Workers Planned: 2"
```

```
Workers Launched: 2"
```

```
-> Parallel Seq Scan on my_tbl (cost=0.00..23311.01 rows=1 width=12)  
(actual time=68.527..101.757 rows=33 loops=3)"
```

```
Filter: ((fld_1 = 500) AND (fld_2 = 100))"
```

```
Rows Removed by Filter: 666634"
```

```
Planning Time: 0.109 ms"
```

```
Execution Time: 116.551 ms"
```

```
SELECT stxname, stxkeys, stxkind FROM pg_statistic_ext;
```

```
>>>
```

stxname	stxkeys	stxkind
my_tbl_fld_1_fld_2_stat	1 2	{d, f, m}

## \* pg\_stat\_advisor does NOT process



- Part-analyzed, TEMP tables;
- UPDATE, DELETE queries;
- Nested Loop, Merge Join, Hash Join nodes
- Columns with ndistinct = 1

## \* TODO

- Kind of statistics
- Remove useless extended statistics
- Adjust default\_statistics\_target


# Thank You for Your Attention



ilia-evdokimov



EvdokimovIlia



ilidock95

