

A large orange graphic consisting of a wavy line that ends in a solid arrow pointing to the right.

pg_stat_advisor - PostgreSQL advisor to create statistics

Ilia Evdokimov, Tantor Labs LLC

Contents



- autovacuum specific
- statistics
- extended statistics
- **pg_stat_advisor**

* autovacuum features

Tables in schema pg_catalog					
Command	Free Space Map (FSM)	Visibility Map (VM)	Heap	Indexes	pg_class
					pg_statistic
VACUUM	Updates	Updates	Removes dead tuples	Removes links to dead tuples	Update attributes: reltuples, relpages, relallvisible
ANALYZE	Not affected	Not affected	Collects usage statistics	Collects usage statistics	Update attributes: reltuples, relpages. But not relallvisible.
					Updates usage statistics for query plan optimization

* GUC-parameters of autovacuum

```
SELECT name, short_desc FROM pg_settings
WHERE name = 'autovacuum_analyze_threshold'
       OR name = 'autovacuum_analyze_scale_factor';
>>>
```

name	short_desc
autovacuum_analyze_scale_factor	Number of tuple inserts, updates, or deletes prior to analyze as a fraction of reltuples.
autovacuum_analyze_threshold	Minimum number of tuple inserts, updates, or deletes prior to analyze.

(2 rows)

* Prepare test table



Create table with disabled `autovacuum` to update statistics manually:

```
CREATE TABLE my_tbl(fld_1 INTEGER, fld_2 BIGINT)
WITH (autovacuum_enabled = false);
>>>
CREATE TABLE
```

* Evaluate data cardinality



Evaluate test data before **INSERT**:

```
SELECT
  COUNT(DISTINCT T.fld_1)::float/10000000,
  COUNT(DISTINCT T.fld_2)::float/10000000
FROM (
  SELECT i/100 AS fld_1, i/500 AS fld_2
  FROM generate_series(1, 10000000) s(i)
) T;
>>>
      0.01      0.002      <-- low cardinality
```

* Insert values

Fill the table with diverse data:

```
INSERT INTO my_tbl (fld_1, fld_2)
SELECT
    i/100 as fld_1,
    i/500 as fld_2
FROM generate_series(1, 10000000) s(i);
>>>
INSERT 0 1000000
```

* Check basic statistics

Show attributes of `my_tbl`:

```
SELECT reltuples, relpages, relallvisible
FROM pg_class
WHERE relname = 'my_tbl';
>>>
```

reltuples	relpages	relallvisible
-1	0	0

(1 row)

* Refresh statistics

Show again attributes of `my_tbl` after `ANALYZE`:

```
ANALYZE VERBOSE my_tbl;
>>>
      INFO:  analyzing "public.my_tbl"

SELECT reltuples, relpages, relallvisible
FROM pg_class
WHERE relname = 'my_tbl';
>>>
 reltuples | relpages | relallvisible
-----+-----+-----
      1e+06 |      5406 |              0
(1 row)
```

* VACUUM

Show attributes of `my_tbl` after `VACUUM`:

```
VACUUM VERBOSE my_tbl;
>>>
INFO:  vacuuming "public.my_tbl"

SELECT reltuples, relpages, relallvisible
FROM pg_class
WHERE relname = 'my_tbl';
>>>
 reltuples | relpages | relallvisible
-----+-----+-----
      1e+06 |      5406 |          5406
(1 row)
```

* default_statistics_target



Show current value:

```
SHOW default_statistics_target;
>>>
  default_statistics_target
-----
          100              <-- this is default value
(1 row)
```

default_statistics_target - sets the default statistics target for table columns without a column-specific target set via **ALTER TABLE SET STATISTICS**. Larger values increase the time needed to do **ANALYZE**, but might improve the quality of the planner's estimates.

* Statistics



```
\x
SELECT *, array_length(histogram_bounds, 1) as histogram_bounds_len
FROM pg_stats
WHERE tablename = 'my_tbl' AND attname = 'fld_1';
>>>
```

```
-[ RECORD 1 ]-----+-----
schemaname      | public
tablename       | tbl
attname         | fld_1
null_frac       | 0
avg_width       | 4
n_distinct      | 1000
most_common_vals | {318,564,596,...}
most_common_freqs | {0.00173333,0.0017,0.00166667,...}
histogram_bounds | {0,8,20,30,39,...}
...
histogram_bounds_len | 101  <-- default_statistics_target
```

* Extended statistics



dependencies	Evaluates functional dependencies between columns
ndistinct	Counts unique combinations of values in columns
mcv	Analyzes the most frequent value combinations in columns

```
CREATE STATISTICS stat_tbl (dependencies)  
ON fld_1, fld_2  
FROM my_tbl;  
>>>  
CREATE STATISTICS
```

* pg_stats_ext



There are no statistics before executing **ANALYZE**:

```
SELECT * FROM pg_stats_ext
WHERE statistics_name = 'stat_tbl';
>>>
      no rows
```

* pg_stats_ext after ANALYZE



```
ANALYZE VERBOSE my_tbl;
>>>
      INFO:  analyzing "public.my_tbl"

SELECT * FROM pg_stats_ext
WHERE statistics_name = 'stat_tbl';
>>>
      "my_tbl"      "{fld_1,fld_2}"      "{""1 => 2": 1.000000}"
```

Remove statistics for the upcoming examples:

```
DROP STATISTICS stat_tbl;
>>>
      DROP STATISTICS
```

* Compare estimated and actual rows

For the best plan **estimated rows** should be similar with **actual rows**:

```
EXPLAIN ANALYZE
SELECT * FROM my_tbl
WHERE fld_1 = 500 AND fld_2 = 100;
>>>
    Gather (cost=1000.00..12656.10 rows=1 width=12) (actual
time=5.148..57.646 rows=100 loops=1)"
    Workers Planned: 2"
    Workers Launched: 2"
    -> Parallel Seq Scan on my_tbl (cost=0.00..11656.00 rows=1
width=12) (actual time=21.646..36.275 rows=33 loops=3)"
        Filter: ((fld_1 = 500) AND (fld_2 = 100))"
        Rows Removed by Filter: 333300"
    Planning Time: 0.128 ms"
    Execution Time: 57.699 ms"
```


* CREATE STATISTICS

Create and collect (refresh) statistics:

```
CREATE STATISTICS stat_tbl (dependencies)
ON fld_1, fld_2
FROM my_tbl;
>>>
CREATE STATISTICS

ANALYZE VERBOSE my_tbl;
>>>
INFO: analyzing "public.my_tbl"
```

* Estimated rows after creating extended statistics



Compare estimated and actual rows after **CREATE STATISTICS**:

```
EXPLAIN ANALYZE
SELECT * FROM my_tbl
WHERE fld_1 = 500 AND fld_2 = 100;
>>>
    Gather  (cost=1000.00..12666.00 rows=100 width=12) (actual
time=5.176..59.361 rows=100 loops=1)"
    Workers Planned: 2"
    Workers Launched: 2"
    -> Parallel Seq Scan on my_tbl  (cost=0.00..11656.00 rows=42
width=12) (actual time=23.620..39.169 rows=33 loops=3)"
        Filter: ((fld_1 = 500) AND (fld_2 = 100))"
        Rows Removed by Filter: 333300"
    Planning Time: 0.165 ms"
    Execution Time: 59.404 ms"
```

* About pg_stat_advisor



What is the Purpose of `pg_stat_advisor`?

- Advises on executing `CREATE STATISTICS` based on specific criteria.
- Suggests performing `ANALYZE` operations in particular scenarios.

Motivation Behind Creating `pg_stat_advisor`

- Faced challenges with complex SQL queries where precise row count estimations significantly impacted query performance.
- Belief in the extension's potential utility for others in the community.

* How use pg_stat_advisor



Manual by command for current session:

```
LOAD 'pg_stat_advisor';  
>>>  
LOAD
```

Or add library to `postgresql.conf`:

```
shared_preload_libraries = 'pg_stat_advisor'
```

Set extension parameters in `postgresql.conf` or current session:

- `pg_stat_advisor.analyze_scale_factor`
- `pg_stat_advisor.suggest_statistics_threshold`

* Continue using of test table



Drop previously created extended statistics and add more data:

```
DROP STATISTICS stat_tbl;
>>>
      DROP STATISTICS

INSERT INTO my_tbl (fld_1, fld_2)
SELECT i/100 AS fld_1, i/500 AS fld_2
FROM generate_series(1000000, 1500000) s(i);
>>>
      INSERT 0 500001
```

* Configure analyze_scale_factor



```
-- n_mod_since_analyze / n_live_tup > pg_stat_advisor.analyze_scale_factor
SET pg_stat_advisor.analyze_scale_factor = 0.01;
>>>
SET

EXPLAIN ANALYZE SELECT * FROM my_tbl WHERE fld_1 = 500 AND fld_2 = 100;
>>>
NOTICE: pg_stat_advisor suggestion: 'ANALYZE my_tbl'
Gather (cost=1000.00..24309.94 rows=1 width=12) (actual
      time=101.366..108.701 rows=100 loops=1)"
  Workers Planned: 2"
  Workers Launched: 2"
    -> Parallel Seq Scan on my_tbl (cost=0.00..23309.84 rows=1
width=12) (actual time=64.040..95.332 rows=33 loops=3)"
      Filter: ((fld_1 = 500) AND (fld_2 = 100))"
      Rows Removed by Filter: 666634"
Planning Time: 0.272 ms"
Execution Time: 114.312 ms"
```

* Configure suggest_statistics_threshold



```
-- estimated_row / actual rows < pg_stat_advisor.suggest_statistics_threshold
SET pg_stat_advisor.suggest_statistics_threshold = 0.01;
>>>
SET
```

```
EXPLAIN ANALYZE SELECT * FROM my_tbl WHERE fld_1 = 500 AND fld_2 = 100;
```

```
>>>
```

```
NOTICE: pg_stat_advisor suggestion: CREATE STATISTICS
my_tbl_fld_1 fld_2 ON fld_1, fld_2 FROM my_tbl
```

```
Gather (cost=1000.00..24311.11 rows=1 width=12) (actual
time=104.973..113.246 rows=100 loops=1)"
```

```
Workers Planned: 2"
```

```
Workers Launched: 2"
```

```
-> Parallel Seq Scan on my_tbl (cost=0.00..23311.01 rows=1 width=12)
(actual time=68.527..101.757 rows=33 loops=3)"
```

```
Filter: ((fld_1 = 500) AND (fld_2 = 100))"
```

```
Rows Removed by Filter: 666634"
```

```
Planning Time: 0.109 ms"
```

```
Execution Time: 116.551 ms"
```

* Results of applying recommendations



```
CREATE STATISTICS stat_tbl (dependencies)
ON fld_1, fld_2
FROM my_tbl;
>>>
CREATE STATISTICS

ANALYZE VERBOSE my_tbl;
>>>
INFO:  analyzing "public.my_tbl"

EXPLAIN ANALYZE
SELECT * FROM my_tbl
WHERE fld_1 = 500 AND fld_2 = 100;
>>>
Gather  (cost=1000.00..24324.71 rows=100 width=12) (actual
        time=102.944..114.427 rows=100 loops=1)"
...
```



Thank You for Your Attention



ilia-evdokimov



EvdokimovIlia



ilidock95

